

Problem Solving with AI Techniques

Stochastic Search

Paul Weng

UM-SJTU Joint Institute

VE593, Fall 2018



JOINT INSTITUTE
交大密西根学院

1 Monte Carlo Tree Search

- Monte Carlo Methods
- Principle of Monte Carlo Tree Search

2 Simulated Annealing

- Local Search
- Principle of Simulated Annealing

Monte Carlo Methods

- Monte Carlo method = sampling from a probability distribution
- It allows to estimate an integral (e.g., \mathbb{E} is an integral)
e.g., If $\mathbb{E}_p[f(X)] = \int_x f(x)p(x)dx$, then

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_i f(x_i) = \mathbb{E}_p[f(X)]$$

Monte Carlo Methods

- Monte Carlo method = sampling from a probability distribution
- It allows to estimate an integral (e.g., \mathbb{E} is an integral)
e.g., If $\mathbb{E}_p[f(X)] = \int_x f(x)p(x)dx$, then

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_i f(x_i) = \mathbb{E}_p[f(X)]$$

or

$$\lim_{N \rightarrow \infty} \sum_i w_i f(x_i) = \mathbb{E}_p[f(X)]$$

- Useful when p is complicated and $\mathbb{E}_p[f(X)]$ cannot be calculated directly

Rejection Sampling

- How can we generate i.i.d. samples $x_i \sim p(x)$?
- **Assumptions:**
 - We can sample $x \sim q(x)$ from a simpler distribution $q(x)$ (e.g., uniform), called **proposal distribution**
 - We can numerically evaluate $p(x)$ for specific x (even if we don't have an analytic expression of $p(x)$)
 - There exists M such that $\forall x, p(x) \leq Mq(x)$ (which implies q has a larger or equal support as p)
- **Rejection Sampling:**
 - Sample a candidate $x \sim q(x)$
 - Accept x with probability $\frac{p(x)}{Mq(x)}$ and reject otherwise
 - Repeat until sample size = N
- This generates an unweighted sample set to approximate $p(x)$

Importance Sampling

- **Assumptions:**

- We can sample $x \sim q(x)$ from a simpler distribution $q(x)$ (e.g., uniform)
- We can numerically evaluate $p(x)$ for specific x (even if we don't have an analytic expression of $p(x)$)

- **Importance Sampling:**

- Sample a candidate $x_i \sim q(x)$
- Add the weighted sample (w_i, x_i) where $w_i = \frac{p(x_i)}{q(x_i)}$
- Repeat N times
- This generates a weighted sample set to approximate $p(x)$
Weights w_i are called importance weights
- **Crucial for efficiency:** a good choice of proposal distribution $q(x)$

- 1 Monte Carlo Tree Search
 - Monte Carlo Methods
 - Principle of Monte Carlo Tree Search

- 2 Simulated Annealing
 - Local Search
 - Principle of Simulated Annealing

Monte Carlo Tree Search (MTCS)

- MCTS is very successful on Computer Go and other games
- MCTS is one of the main components in Alphago
- MCTS is a quite novel technique (~ 10 years)
- MCTS is rather simple to implement
- MCTS is very general: applicable on any discrete domain

Flat Monte Carlo


- Goal of MCTS: estimate the expected value of action (*Q-function*)

$$Q(s, a) = \mathbb{E}[\Delta \mid s, a]$$

where expectation is taken over future randomized actions (including possible adversary choices and possible stochastic transitions in the environment)

- In search trees, MCTS provides an estimate of f
- *Flat Monte Carlo* does so by rolling out many random simulations (using a `ROLLOUTPOLICY`) without growing a tree
- Key difference/advantage of MCTS over flat Monte Carlo: search focuses computational effort on promising actions

Generic MCTS Scheme



from Browne et al.

- 1: start tree $V = \{v_0\}$
- 2: **while** within computational budget **do**
- 3: $v_l \leftarrow \text{TREEPOLICY}(V)$ chooses a leaf of V
- 4: append v_l to V
- 5: $\Delta \leftarrow \text{ROLLOUTPOLICY}(V)$ rolls out a full simulation, with return Δ
- 6: $\text{BACKUP}(v_l, \Delta)$ updates the values of all parents of v_l
- 7: **end while**
- 8: return best child of v_0

Generic MCTS Scheme: Remarks

- Like flat MC, MCTS typically computes full roll-outs to a terminal state. A heuristic to estimate the utility of a state is not needed, but can be incorporated.
- The tree grows unbalanced.
- TREEPOLICY decides where the tree is expanded – and needs to trade-off exploration vs. exploitation.
- ROLLOUTPOLICY is necessary to simulate a roll-out. It typically is a random policy (at least a randomized policy).

Upper Confidence Tree (UCT)

- UCT uses UCB to realize TREEPOLICY, i.e., to decide where to expand the tree
- BACKUP updates all parents of v_l as
$$n(v) \leftarrow n(v) + 1 \text{ (count how often it has been tried)}$$
$$Q(v) \leftarrow Q(v) + \Delta \text{ (sum of rewards received)}$$
- TREEPOLICY chooses child nodes based on UCB:

$$\arg \max \frac{Q(v')}{n(v')} + \beta \sqrt{\frac{2 \ln n(v)}{n(v')}}}$$

or chooses v' if $n(v') = 0$

1 Monte Carlo Tree Search

- Monte Carlo Methods
- Principle of Monte Carlo Tree Search

2 Simulated Annealing

- Local Search
- Principle of Simulated Annealing


Iterative Improvement Algorithms

- In some problems, *path* is irrelevant, *goal state* is the solution

$$\begin{array}{rcccccc}
 & & S & E & N & D \\
 + & & M & O & R & E \\
 \hline
 = & M & O & N & E & Y
 \end{array}$$


- State space = set of *complete* configurations
 find configuration satisfying all constraints, e.g., timetable
 find *optimal* configuration, e.g., Travelling Salesperson Problem
- Iterative improvement algorithms:** keep a single "current" state, try to improve it
- Constant space complexity

Example: n -Queens




- **Goal:** Place n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- **Algorithmic principle:**
 - Start with n queens placed on board
 - Repeatedly, move a queen to reduce number of conflicts

Example: n -Queens




- **Goal:** Place n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- **Algorithmic principle:**
 - Start with n queens placed on board
 - Repeatedly, move a queen to reduce number of conflicts
- Can solve n -queens problems almost instantaneously for very large n , e.g., $n = 10^6$

Example: Travelling Salesperson Problem



- **Goal:** Find minimal-cost cycle
- **Algorithmic principle:**
 - Start with any complete tour
 - Repeatedly, perform pairwise exchanges

Example: Travelling Salesperson Problem



- **Goal:** Find minimal-cost cycle
- **Algorithmic principle:**
 - Start with any complete tour
 - Repeatedly, perform pairwise exchanges
- Variants of this approach get within 1% of optimal very quickly with thousands of cities

Hill-climbing (or gradient ascent/descent)


”Like climbing Everest in thick fog with amnesia”
(assuming we maximize an utility function)

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                   neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```


Hill-climbing (contd.)

- Hill-climbing can get stuck, see **state space landscape**



Hill-climbing (contd.)


- Hill-climbing can get stuck, see **state space landscape**



- Random-restart hill climbing:**
overcomes local maxima — trivially complete

Hill-climbing (contd.)

- Hill-climbing can get stuck, see **state space landscape**



- Random-restart hill climbing:**
overcomes local maxima — trivially complete
- Random sideways moves:**
 \oplus escape from shoulders, \ominus loop on flat maxima

1 Monte Carlo Tree Search

- Monte Carlo Methods
- Principle of Monte Carlo Tree Search

2 Simulated Annealing

- Local Search
- Principle of Simulated Annealing

Simulated Annealing

- **Idea:** Escape local maxima by allowing some “bad” moves
but gradually decrease their size and frequency

function **SIMULATED-ANNEALING**(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

local variables: *current*, a node

next, a node

T, a “temperature” controlling prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] – VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E / T}$

Properties of simulated annealing

- At fixed “temperature” T , state occupation probability reaches Boltzman distribution

$$p(x) \propto e^{-\frac{E(x)}{kT}}$$

T decreased slowly enough \implies always reach best state x^*

because $e^{-\frac{E(x^*)}{kT}} / e^{-\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \gg 1$ for small T

- Is this necessarily an interesting guarantee?
- Devised by Metropolis et al., 1953, for physical process modelling
- Widely used in VLSI layout, airline scheduling, etc.

Going Further...

- **Local beam search:** keep k states instead of 1; choose top k of all their successors
- **Genetic algorithms:** stochastic local beam search + generate successors from *pairs* of states
- **Evolutionary Strategies** (e.g., CMA-ES)
- **Nature-inspired algorithms** (e.g., Particle Swarm Optimization)
- **Metaheuristics** (e.g., Cross-Entropy Method)