

Problem Solving with AI Techniques

Sequential Decision-Making under Uncertainty

Paul Weng

UM-SJTU Joint Institute

VE593, Fall 2018




JOINT INSTITUTE
交大密西根学院

- 1 Introduction
- 2 Markov Decision Process
- 3 Dynamic Programming


Decision-making

- **Goal:** design adaptive autonomous systems that can act on our behalf
- **Examples:**
 - Robotics
 - Traffic light controller
 - Data center control
 - Streaming video player
 - Intelligent tutoring system
- Interaction loop



Example: Autonomous Navigation

- **Goal:** move autonomously from point A to point B




- Repeat: choose an action, observe new state, check if goal state
- Optimize travel time, power consumption...
- Search problem if effects of actions are deterministic
- Stochastic model: exogenous events, imperfect sensors/actuators, model error...

- 1 Introduction
- 2 Markov Decision Process**
- 3 Dynamic Programming

Markov Decision Process

Markov Decision Process (MDP) defined as a quadruplet $(\mathcal{S}, \mathcal{A}, T, R)$

- \mathcal{S} is a set of states
- \mathcal{A} is a set of actions
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ is a transition function, $T(s, a, s') = \mathbb{P}(s' | s, a)$
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ is a reward function



Examples of MDP

An MDP is defined by:

- \mathcal{S} set of states
- \mathcal{A} set of actions
- $T(s, a, s')$ transition function
- $R(s, a)$ reward function

Examples of MDP

An MDP is defined by:

- \mathcal{S} set of states
- \mathcal{A} set of actions
- $T(s, a, s')$ transition function
- $R(s, a)$ reward function

Example: Navigation of a robot

- Positions of the robot
- Possible moves
- Probabilistic effects of a move
- Reward encoding a certain task

Examples of MDP

An MDP is defined by:

- \mathcal{S} set of states
- \mathcal{A} set of actions
- $T(s, a, s')$ transition function
- $R(s, a)$ reward function

Example: Inventory control

- Stock levels
- Orders
- Stochastic demand
- Earnings - costs

Examples of MDP

An MDP is defined by:

- \mathcal{S} set of states
- \mathcal{A} set of actions
- $T(s, a, s')$ transition function
- $R(s, a)$ reward function

Example: Video games

- Resources, Units, Structures
- Explore, Exploit, Build...
- Uncertain consequences
- Utility

Examples of MDP

An MDP is defined by:

- \mathcal{S} set of states
- \mathcal{A} set of actions
- $T(s, a, s')$ transition function
- $R(s, a)$ reward function

Generally, T and R are not known \Rightarrow Reinforcement learning

Moreover, if there is only one state with random rewards \Rightarrow Multi-armed bandit

Goal in an MDP

- **High-level goal:** determine which action to choose in every state at every time step
- **Important notions:**
 - Trajectory s_0, s_1, s_2, \dots : sequence of states
 - History $s_0, a_0, s_1, a_1, s_2, \dots$: sequence of state/action
 - Policy π : stationary or not, Markov/history-dependent, deterministic/randomized
 - Horizon: finite or infinite
- **Problem:** find "best" policy for a given MDP

Discussions

- **Assumptions**
 - Stationary (time-homogeneous) model
 - Markov property
 - Discrete time
 - States are observable (MDP can be generalized to Partially Observable MDP or POMDP)
- MDP can be seen as an extension of
 - a deterministic state-space where actions have stochastic effects
 - a Markov chain where in every state we can choose the distribution over next states

How to Define "Best"?

- Value of a history = sum of rewards (possibly discounted by $\gamma \in (0, 1)$)
- A policy induces a probability distribution over histories
- Value $v^\pi(s)$ of a policy in a state = expectation of the value of histories it can generate
- This defines the value function v^π of a policy π
- Optimal value function $v^*(s) = \max_\pi \mathbb{E}_\pi[\sum_t \gamma^t R(S_t, A_t) \mid S_0 = s]$
- **Theorem:** For **finite** horizon problems, finite MDP admits an optimal policy that is stationary, Markov, and deterministic
- **Theorem:** For **infinite** horizon problems, finite MDP admits an optimal policy that is stationary, Markov, and deterministic

- 1 Introduction
- 2 Markov Decision Process
- 3 Dynamic Programming**

Induction for Value Function of Fixed Policy

- Induction for π :

$$v_0^\pi(s) = 0$$

$$v_t^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') v_{t-1}^\pi(s')$$

- v_T^π provides the value function of π for finite horizon T
- v_t^π converges to value function v^π of π for infinite horizon
- value function of π for infinite horizon satisfies:

$$v^\pi(s) = R(s, \pi) + \gamma \sum_{s'} T(s, \pi, s') v^\pi(s')$$

Backward Induction and Value Iteration

- Induction:

$$v_0^*(s) = 0$$

$$v_t^*(s) = \max_a R(s, a) + \gamma \sum_{s'} T(s, a, s') v_{t-1}^*(s')$$

- v_T^* provides optimal value function for finite horizon T
- v_t^* converges to optimal value function v^* for infinite horizon
- Optimal value function for infinite horizon satisfies:

$$v^*(s) = \max_a R(s, a) + \gamma \sum_{s'} T(s, a, s') v^*(s')$$

Value Iteration: Algorithm

```

1 ValueIteration(MDP,  $\varepsilon$ )
2  $\forall s, v^*(s) = 0$ 
3 repeat
4   for  $s \in \mathcal{S}$  do
5     for  $a \in \mathcal{A}$  do
6        $Q^*(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') v^*(s')$ 
7        $nv(s) \leftarrow \max_a Q^*(s, a)$ 
8    $v^* \leftarrow nv$ 
9 until  $\|v^* - nv\| \leq \varepsilon$ ;
0 return  $Q^*$ 
  
```

- Computational complexity for one step: $O(|\mathcal{S}|^2|\mathcal{A}|)$
- Computational complexity: $O(\text{poly}(|\mathcal{S}|, |\mathcal{A}|, \frac{1}{1-\gamma}))$

Policy Iteration: Principle

- Given current policy π , improve it with one-step look-ahead
- Alternate between
 - Policy evaluation of current policy
$$v_t^\pi(s) = \max_a R(s, a) + \gamma \sum_{s'} T(s, a, s') v_{t-1}^\pi(s')$$
 - Policy improvement of current policy:
$$\pi'(s) = \arg \max_a R(s, a) + \gamma \sum_{s'} T(s, a, s') v^\pi(s')$$

Policy Iteration

```

1 PolicyIteration(MDP)
2   initialize  $\pi$ 
3   repeat
4      $\pi' \leftarrow \pi$ 
5     compute  $v^\pi$  (possibly approximately)
6     for  $s \in \mathcal{S}$  do
7       for  $a \in \mathcal{A}$  do
8          $Q^\pi(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') v^\pi(s')$ 
9          $\pi(s) \leftarrow \arg \max_a Q^\pi(s, a)$ 
10  until  $\pi \neq \pi'$ ;
11  return  $\pi$ 
  
```

- Computational complexity: $O(\text{poly}(|\mathcal{S}|, |\mathcal{A}|, \frac{1}{1-\gamma}))$