

Problem Solving with AI Techniques

Reinforcement Learning

Paul Weng

UM-SJTU Joint Institute

VE593, Fall 2018




JOINT INSTITUTE
交大密西根学院

- 1 Overview
- 2 Algorithms

What is RL?

- General framework for learning from interactions



- Inspired by animal/human learning
 - dopamine-based learning
 - animal training
- MDP, but unknown model
- Contextual MAB with action-independent transition

Applications

- Video games, e.g., [breakout](#)
- Board game, e.g.,
Backgammon, Go, Chess, Shogi
- Robotic control



Types of Problems

- Infinite horizon
 - Ergodic MDP
- Repeated problems:
 - Finite horizon: episodic problem
 - Goal oriented problem

Two families of approaches:

- Model-based RL
 - Indirect approach

- Model-free RL
 - Value-based approaches
 - Policy search

1 Overview

2 Algorithms

- Model-Based RL
- Value-Based Approaches
- Linear Approximation
- Policy Search

Model-based RL

- Learn model first
 - Estimate $T(s, a, \cdot)$ for all s, a
 - Estimate $R(s, a)$ for all s, a
- Find optimal policy using learned model
- **Issue:** Hard to learn model (in terms of sample, space requirement...)
- As learned model is an approximation, optimal policy in learned model is probably not optimal in true environment
- Maybe a good approach for multi-task RL


1 Overview

2 Algorithms

- Model-Based RL
- Value-Based Approaches
- Linear Approximation
- Policy Search

Monte Carlo Sampling: Principle

- **Goal:** $\max_{\pi} \mathbb{E}_{\pi} [\sum_t \gamma^t R(S_t, A_t) \mid S_0 = s]$
- **Monte Carlo approach:** Sample many trajectories



- MCTS in expectimax tree

Monte Carlo Sampling: Incremental Estimation

- Assume we can sample $n + 1$ times from $T(s, \pi(s), \cdot)$

$$\begin{aligned}\hat{v}_{n+1}^{\pi}(s) &= \frac{v_1 + \dots + v_{n+1}}{n + 1} \\ &= \frac{n}{n + 1} \hat{v}_n^{\pi}(s) + \frac{1}{n + 1} v_{n+1} \\ &= \hat{v}_n^{\pi}(s) + \frac{1}{n + 1} (v_{n+1} - \hat{v}_n^{\pi}(s)) \\ &\sim \hat{v}_n^{\pi}(s) + \alpha (v_{n+1} - \hat{v}_n^{\pi}(s))\end{aligned}$$

- Property:** no bias, but high variance
- Issue:** May be inefficient, as it doesn't exploit the problem structure

Boostrapping

- Recall

$$\begin{aligned}
 v^\pi(s) &= \mathbb{E}_\pi \left[\sum_t \gamma^t R(S_t, A_t) \mid S_0 = s \right] \\
 &= \mathbb{E} \left[\sum_t \gamma^t R(S_t, A_t) \mid S_0 = s, \pi(S_t) = A_t, S_{t+1} \sim T(S_t, A_t, \cdot) \right] \\
 &= R(s, \pi(s)) + \gamma \mathbb{E}_{S' \sim T(s, \pi(s), \cdot)} [v^\pi(S')] \\
 v^*(s) &= \max_\pi \mathbb{E}_\pi \left[\sum_t \gamma^t R(S_t, A_t) \mid S_0 = s \right] \\
 &= \max_a R(s, a) + \gamma \mathbb{E}_{S' \sim T(s, a, \cdot)} [v^*(S')]
 \end{aligned}$$

- Idea:** Use current estimates to avoid sampling whole trajectories

$$\begin{aligned}
 \hat{v}^\pi(s) &\leftarrow R(s, \pi(s)) + \gamma \mathbb{E}_{S' \sim T(s, \pi(s), \cdot)} [\hat{v}^\pi(S')] \\
 \hat{v}^*(s) &\leftarrow \max_a R(s, a) + \gamma \mathbb{E}_{S' \sim T(s, a, \cdot)} [\hat{v}^*(S')]
 \end{aligned}$$

Temporal Difference

- Incremental updates:

$$\hat{v}_{n+1}^{\pi}(s) = \hat{v}_n^{\pi}(s) + \alpha(r_{n+1} + \gamma \hat{v}_n^{\pi}(s') - \hat{v}_n^{\pi}(s))$$

- TD error: $r_{n+1} + \gamma \hat{v}_n^{\pi}(s') - \hat{v}_n^{\pi}(s)$
- TD(0): learns the value function of a policy
- can be extended to the Q function:

$$\hat{Q}_{n+1}^{\pi}(s, a) = \hat{Q}_n^{\pi}(s, a) + \alpha(r_{n+1} + \gamma \hat{Q}_n^{\pi}(s', a') - \hat{Q}_n^{\pi}(s, a))$$

- Why use the Q function?
- Property:** biased, but low variance

Sarsa: Algorithm

- **Idea:** Estimate $Q^\pi(s, a)$ for all s, a and improve π

```
1 initialize  $Q(s, a)$ 
2 for each episode do
3   initialize  $s$ 
4   choose  $a$  in  $s$  using  $Q$  and possibly some randomness
5   repeat
6     observe  $s', r$  after applying  $a$  in  $s$ 
7     choose  $a'$  in  $s'$  using  $Q$  and possibly some randomness
8      $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$ 
9      $s \leftarrow s'; a \leftarrow a'$ 
10  until  $s$  is terminal;
```

Exploration Policy

- **Greedy policy:** $\pi(s) = \arg \max_a Q(s, a)$
- **ϵ -greedy policy:** choose $\arg \max_a Q(s, a)$ with probability $1 - \epsilon$ and a random action otherwise
- **softmax policy:** choose a with probability
$$\frac{\exp \frac{Q(s,a)}{\tau}}{\sum_{a'} \exp \frac{Q(s,a)}{\tau}}$$

Sarsa: Convergence

- Convergence to optimal policy and optimal Q-function if:
 - π converges to the greedy policy
 - But, all pairs s, a are visited infinitely often

- **Example:** choose π as $\frac{1}{t}$ -greedy policy

Q-learning: Algorithm

- **Idea:** Estimate $Q^*(s, a)$ for all s, a

```
1 initialize  $Q(s, a)$ 
2 for each episode do
3   initialize  $s$ 
4   repeat
5     choose  $a$  in  $s$  using  $Q$  and possibly some randomness
6     observe  $s', r$  after applying  $a$  in  $s$ 
7      $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$ 
8      $s \leftarrow s'$ 
9   until  $s$  is terminal;
```

Q-learning: Convergence

- Convergence to optimal policy and optimal Q-function if:
 - all pairs s, a are visited infinitely often
 - learning rates satisfy: $\forall s, a$



$$\sum_t \alpha_t(s, a) = \infty \quad \sum_t \alpha_t^2(s, a) < \infty$$

- In practice, $\forall s, a, t, \alpha_t(s, a) = \alpha$, which helps in non-stationary environments

Sarsa vs Q-learning

- **Behavior policy:** policy used in environment
- **Target policy:** policy whose value (or Q) function is learned
- Off-policy vs on-policy algorithms
- Online performance is preferred \Rightarrow Sarsa
- Learning optimal Q-function is important, no need to control behavior policy \Rightarrow Q-learning

Example: Cliff Walking



from Sutton and Barto

Issues

- **Limitation:** These methods assume finite MDPs
- **Curse of dimensionality**
- Need of compact representations
- We'll focus on methods based on function approximation

1 Overview

2 Algorithms

- Model-Based RL
- Value-Based Approaches
- **Linear Approximation**
- Policy Search

Linear Approximation for Value Functions

- Assume we want to approximate $v^\pi : \mathcal{S} \rightarrow \mathcal{A}$
- Basis Functions:** $\phi = (\phi_1, \phi_2, \dots, \phi_K)^\top$ where $\phi_k : \mathcal{S} \rightarrow \mathbb{R}$
- Goal:** We want to find \mathbf{w} such that $v_{\mathbf{w}} = \mathbf{w}^\top \phi$ is close to v^π
- e.g., $\arg \min_{\mathbf{w}} \mathbb{E}_\mu[(v^\pi(S) - v_{\mathbf{w}}(S))^2 \mid S \sim \mu]$
- Idea:** Optimize its empirical version $\frac{1}{T} \sum_t (v_{\mathbf{w}}(s_t) - v^\pi(s_t))^2$ with stochastic gradient descent:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \nabla_{\mathbf{w}} v_{\mathbf{w}}(s_t) (v^\pi(s_t) - v_{\mathbf{w}}(s_t))$$

- Issue:** v^π not known \Rightarrow bootstrapping
- Linear TD(0):** $\mathbf{w} \leftarrow \mathbf{w} + \alpha \phi(s_t) (r_t + \gamma \mathbf{w}^\top \phi(s_{t+1}) - \mathbf{w}^\top \phi(s_t))$

Linear Approximation for Q-Functions

- Previous approach can be extended to Q-functions
- **Basis Functions:** $\phi = (\phi_1, \phi_2, \dots, \phi_K)^\top$ where $\phi_k : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

- **Linear Sarsa:**

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \phi(s_t, a_t) (r_t + \gamma \mathbf{w}^\top \phi(s_{t+1}, a_{t+1}) - \mathbf{w}^\top \phi(s_t, a_t))$$

- **Linear Q-learning:**

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \phi(s_t, a_t) (r_t + \gamma \max_a \mathbf{w}^\top \phi(s_{t+1}, a) - \mathbf{w}^\top \phi(s_t, a_t))$$

Discussions

- Strictly speaking, previous method is not a gradient method, because of unknown target
- Other formulations of objective functions are possible, see Sutton and Barto's book
- **Issue:** How to define the basis functions in ϕ ?
- Previous approach could be used with non-linear function approximators, however it may be unstable
- What if the action space is continuous?

1 Overview

2 Algorithms

- Model-Based RL
- Value-Based Approaches
- Linear Approximation
- Policy Search

Policy Search

- **Issue:** Previous methods do not work with continuous action space
- **Idea:** Try to find a good (possibly stochastic) policy directly
- **Possible methods:** Local (stochastic) search, metaheuristics...
- **Policy gradient:** gradient descent/ascent-based method
- All those methods may converge to a local optimum

Policy Gradient

- **Assumptions:** Differentiable parametrized policy $\pi_\theta(s, a)$
- **Objective function:**

$$J(\theta) = \mathbb{E}\left[\sum_t \gamma^t R(S_t, A_t) | S_0 \sim \mu, A_t \sim \pi_\theta(S_t, \cdot), S_{t+1} \sim T(S_t, A_t, \cdot)\right]$$

- **Gradient update:** $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
- **Likelihood trick:** $\nabla_\theta \pi_\theta(s, a) = \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)$
- **Policy Gradient Theorem:** $\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) Q^{\pi_\theta}(s, a)]$
- **Reinforce:** Approximate the expectations by sampling

Examples of Parametrized Policy

- Softmax policy

- Define feature functions $\phi(s, a) = (\phi_1(s, a), \dots, \phi_K(s, a))$
- $\pi_{\theta}(s, a) \propto e^{\phi(s, a)^{\top} \theta}$
- $\nabla_{\theta} \log \pi_{\theta}(s, a) = \phi(s, a) - \mathbb{E}_{\pi_{\theta}}[\phi(s, \cdot)]$

- Gaussian policy

- Define feature functions $\phi(s) = (\phi_1(s), \dots, \phi_K(s))$
- $\pi_{\theta}(s, \cdot) = \mathcal{N}(\phi(s)^{\top} \theta, \sigma^2)$
- $\nabla_{\theta} \log \pi_{\theta}(s, a) = \frac{(a - \phi(s)^{\top} \theta) \phi(s)}{\sigma^2}$